

Laboratory project report M1
Parallel programming for neural networks

Djahid ABDELMOUMENE
Encadrant technique: Philippe GAUSSIER
Université de Cergy-Pontoise ETIS

January 11, 2020

Contents

1	Introduction	3
2	optimisation techniques	3
2.1	Naive algorithm	3
2.2	Loop nest optimisations	4
2.3	Parallelization	6
2.4	Parallelization and optimisations	7
3	Algorithm comparison	9
4	Code analysis	10
5	Cause of the problem	11
5.1	The cache memory	11
5.1.1	DRAM vs SRAM	12
5.1.2	Cache lines	13
5.1.3	The cause	14
6	Solution and results	15
6.1	tests	17
7	Conclusion	18

1 Introduction

As processors get faster and faster than memory access times, optimisations start to become necessary in order to utilize the CPU to its full capacity, this is especially true when trying to write parallel code, since memory access bottlenecks can throttle the performance greatly, because the access to memory can't be parallelized which causes race conditions to occur, rendering the parallel code useless. To avoid this problem the parallel code must utilize the cache or some sort of secondary SRAM that allows faster access times at the cost of its storage capacity, which is usually tolerable using some techniques that will be discussed later on.

The main problem addressed here is the optimisation of a neural network simulator called 'promethe'. the main objective is to use parallel programming to scale up the performance of the simulator, as well as solve memory access bottlenecks that are hindering the parallelization of the code.

2 optimisation techniques

In this section we review some optimisation techniques that are used to increase the locality of memory use, that is, increase the reuse of values loaded into the cache. We'll also see the syntax of OpenMP parallelization and its effects on the performance.

To view these techniques well try to optimize matrix multiplication of two $N \times N$ matrices in C, We will be sticking to the naive algorithm with a complexity of $\mathcal{O}(n^3)$.

2.1 Naive algorithm

As a baseline for the performance comparison we'll take the naive version of the matrix multiplication as shown here:

```
1 for(i=0; i<N; i++) {
2     for(j=0; j<N; j++) {
3         acc = 0.0;
4         for(x=0; x<N; x++) {
5             acc += a[i][x] * b[x][j];
6         }
7         res[i][j] = acc;
8     }
9 }
```

And here we can see the benchmark results of the algorithm with different levels of compiler optimisation levels (Here we used GCC v9.2), you should note

that all the tests here will be on a machine with an intel 7-6700HQ CPU running at 2.60GHz with 4 cores.

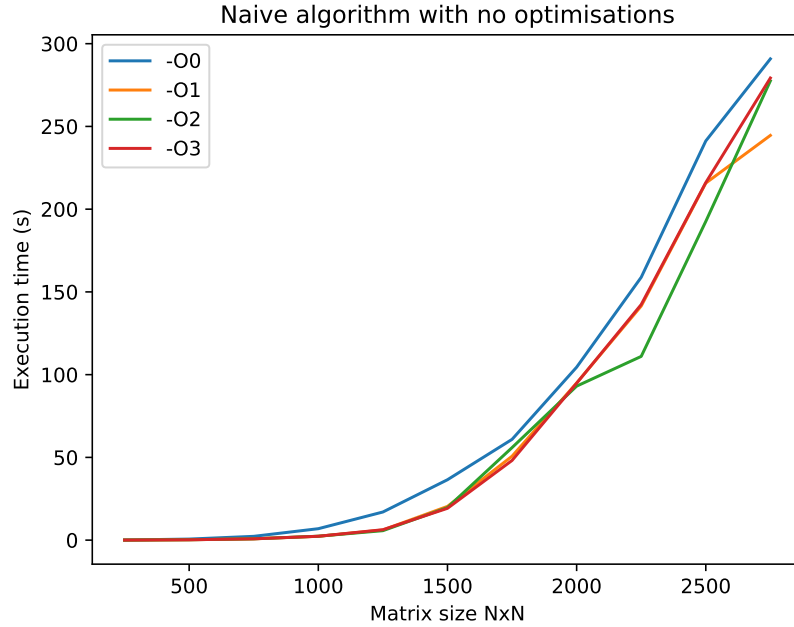


Figure 1: Naive algorithm performance results

2.2 Loop nest optimisations

Here we have a more optimized version [1] [2] [3], with many optimisation techniques used:

```

1 int ib = 10, kb = 10;
2 for (ii = 0; ii < N; ii += ib) {
3     for (kk = 0; kk < N; kk += kb) {
4         for (j=0; j < N; j += 2) {
5             for(i = ii; i < ii + ib; i += 2 ) {
6                 if (kk == 0)
7                     acc00 = acc01 = acc10 = acc11 = 0;
8                 else {
9                     acc00 = res[i + 0][j + 0];
10                    acc01 = res[i + 0][j + 1];
11                    acc10 = res[i + 1][j + 0];
12                    acc11 = res[i + 1][j + 1];
13                }
14                for (k = kk; k < kk + kb; k++) {
15                    acc00 += b[k][j + 0] * a[i + 0][k];

```

```

16         acc01 += b[k][j + 1] * a[i + 0][k];
17         acc10 += b[k][j + 0] * a[i + 1][k];
18         acc11 += b[k][j + 1] * a[i + 1][k];
19     }
20     res[i + 0][j + 0] = acc00;
21     res[i + 0][j + 1] = acc01;
22     res[i + 1][j + 0] = acc10;
23     res[i + 1][j + 1] = acc11;
24 }
25 }
26 }
27 }

```

The first technique used here is register level tiling, which is why we see four different accumulators, and by doing this we reuse every loaded value from matrices a and b twice. which reduces the amount of memory loads necessary.

Next we have the i loop blocking. And by blocking we mean the addition of a second loop (ii loop) that takes larger steps (ib) while keeping the original loop to go through the gaps left by the bigger loop. this technique is used to increase the locality of the data since we are parsing each matrix in rectangular strips. We also have the k loop blocked by a factor (kb) which means we'll be parsing the matrices in square strips of sizes $ibxkb$, this will have a major impact on the performance since each square will be reused multiple times, which reduces the amount of memory loads necessary significantly. These loop blocks require some further treatment in order for it to function correctly, for example here we need to explicitly tell it to initialize the accumulators at $kk == 0$ because we aren't viewing the values contiguously. There are some more special cases that have been omitted from here that make the code work for values of N that aren't multiples of ib and kb .

And here we can see some major performance improvements over the naive version. as well as the fact that these techniques depend heavily on the eventual compiler's optimisations, because the performance spikes on when the compiler optimisations are turned off (optimisation level -O0).

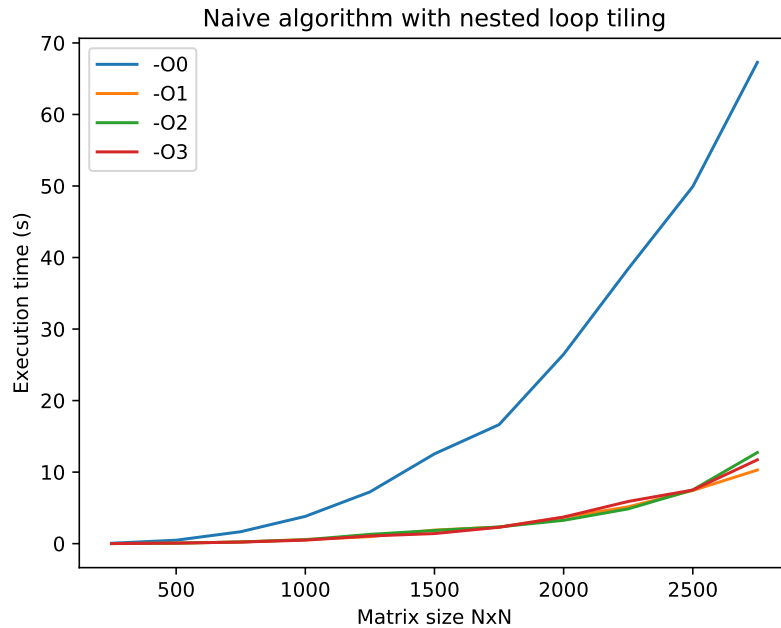


Figure 2: with register tiling and loop blocking

2.3 Parallelization

Here we can see the OpenMP syntax, and we can notice that OpenMP works with C's pragma directives, the 'omp' part is to specify the OpenMP directives, and 'parallel for' is to say that the next for loop needs to be run in parallel, next we specify the variables that should be either shared or made private, to avoid data racing, and we also have the scheduling method for the parallelization, which in this case is static, meaning the loop's iterations will be distributed uniformly between the threads.

```

1 #pragma omp parallel for shared(a, b, res) \
2   private(i, j, x, acc) \
3   schedule(static)
4 for(i=0; i<N; i++) {
5     for(j=0; j<N; j++) {
6         acc = 0.0;
7         for(x=0; x<N; x++) {
8             acc += a[i][x] * b[x][j];
9         }
10        res[i][j] = acc;
11    }
12 }
```

And as visible in the graph, there is a 4 time increase in performance as expected from a 4 core machine.

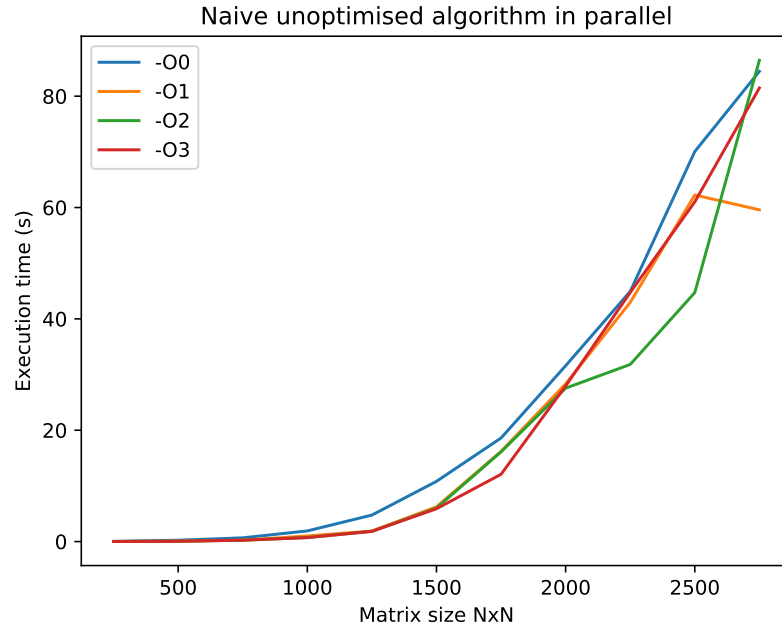


Figure 3: With OpenMP parallelization

2.4 Parallelization and optimisations

Here we have both the parallel part and optimisation techniques.

```

1 int ib = 10, kb = 10;
2 #pragma omp parallel for shared(a, b, res, ib, kb) \
3   private(i, ii, j, k, kk, acc00, acc01, acc10, acc11) \
4   schedule(static)
5 for (ii = 0; ii < N; ii += ib) {
6   for (kk = 0; kk < N; kk += kb) {
7     for (j=0; j < N; j += 2) {
8       for(i = ii; i < ii + ib; i += 2 ) {
9         if (kk == 0)
10            acc00 = acc01 = acc10 = acc11 = 0;
11         else {
12            acc00 = res[i + 0][j + 0];
13            acc01 = res[i + 0][j + 1];
14            acc10 = res[i + 1][j + 0];
15            acc11 = res[i + 1][j + 1];
16        }

```

```

17         for (k = kk; k < kk + kb; k++) {
18             acc00 += b[k][j + 0] * a[i + 0][k];
19             acc01 += b[k][j + 1] * a[i + 0][k];
20             acc10 += b[k][j + 0] * a[i + 1][k];
21             acc11 += b[k][j + 1] * a[i + 1][k];
22         }
23         res[i + 0][j + 0] = acc00;
24         res[i + 0][j + 1] = acc01;
25         res[i + 1][j + 0] = acc10;
26         res[i + 1][j + 1] = acc11;
27     }
28 }
29 }
30 }

```

And we can see that the optimisation techniques scale up with the parallelization, which implies that there are no bottlenecks.

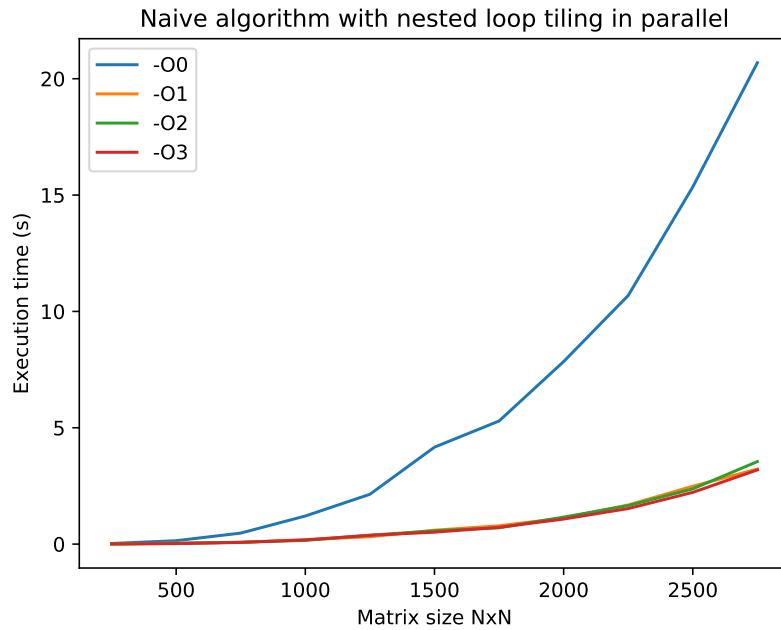


Figure 4: With the optimisation and in parallel

3 Algorithm comparison

Here we compare the four different programs on different compiler optimisation levels.

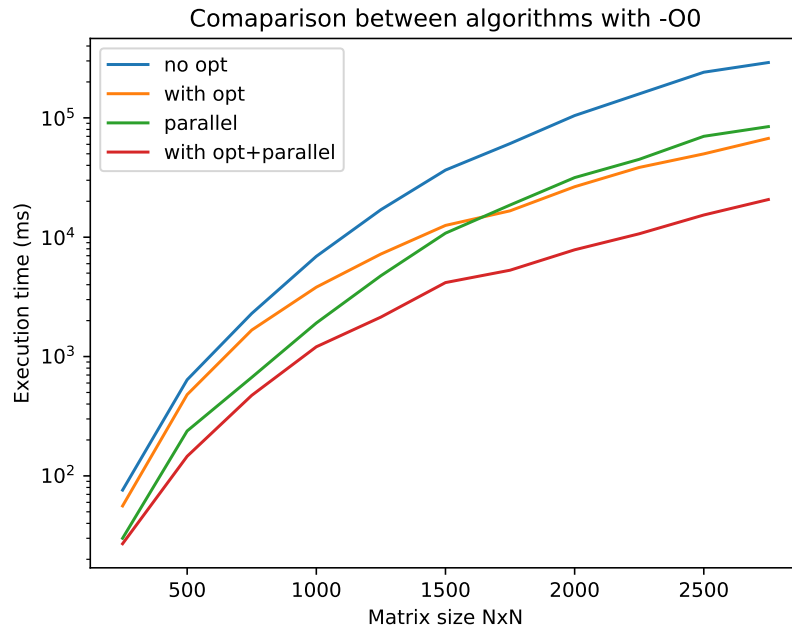


Figure 5: comparison with GCC -O0 (on a logarithmic scale)

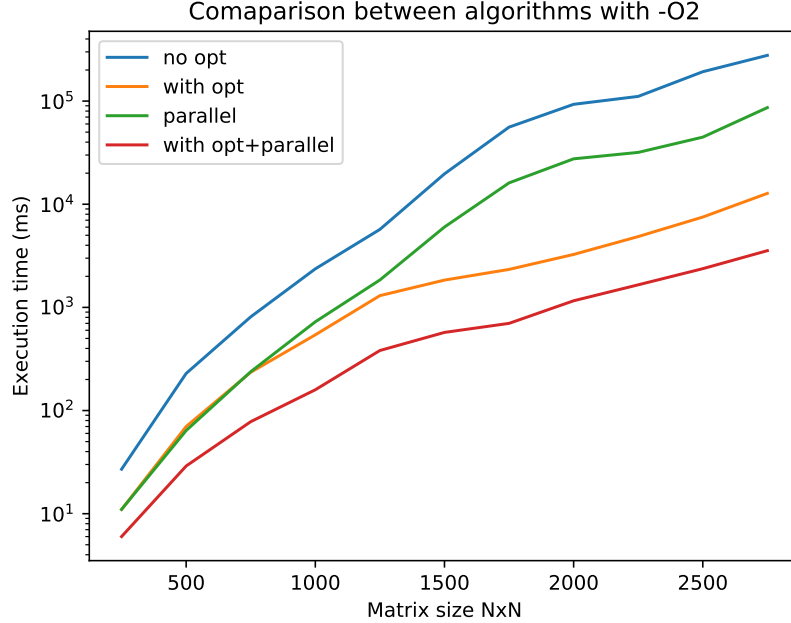


Figure 6: comparison with GCC -O2 (on a logarithmic scale)

We can notice that the parallel naive version is less affected by the compiler optimisation level, and the loop tiling has a better performance than just direct parallelization of the code, and of course the combination of the two gives a major performance improvement.

4 Code analysis

With the goal of optimizing the neural network simulator promethe, I started by trying to parallelize parts of the code using OpenMP, after setting up the necessary compiler flags and libraries for OpenMP.

To benchmark the performance I set up a simple simulation with a hebbian NN and a noise function as an input, along with a function to measure the elapsed time of each update (`f_display_framerate`).

The first bit of code was in `prom_user/src/main.c` where I set up OpenMP to run the procedure responsible for launching the update functions of the neural networks in parallel, and here the performance didn't scale up with the number of threads (with 4 threads 20% improvement).

After thorough examination of the function responsible for updating the neurons's coefficients, which performed a dot product of the neurons coefficients with their outputs. The problem appeared to come from the memory access to the neurons coefficients, after a simple test using a simpler 2d array as a replacement for the coefficients' values. the problem seemed to be fixed along with a much better execution time.

To analyze the cause of this problem I did some research on the subject where it turned out that the problem was caused by the size of the coefficients structure (`type_coeff`) which was 56 bytes long, this size made it so that each memory access operation had to be retrieved from the main memory (as opposed to the cache) at each request from the inner (`calcule_produit`) loop. that is, the entire values were loaded into the cache because the structure took up almost all of the cache line (usually 64 bytes), that meant that most memory accesses missed the cache since only about only one element was loaded from the coefficients array every time. which caused a bottleneck that encumbered the parallelization because each thread has to wait for the other threads' memory requests.

5 Cause of the problem

To track down the source of the problem, that is, what is causing the coefficient's structure (`type_coeff`) to throttle the performance. We are going to see basic a overview of how the cache and cache lines work, and how data is loaded into it.

5.1 The cache memory

The cache is a hardware component usually located in the CPU that allows much faster access times to memory at the expense of the storage size.

The main different between the cache and the main memory are the cells that construct them, where the cache cells are called SRAM, as opposed to the DRAM cells in the main memory.

In modern CPUs there are multiple levels (typically 3, L1, L2 and L3) of cache where L1 is a few tens of KBs, and L2 a few hundreds of KBs, and the third level a few MBs, and the first level is usually much faster than the second which in turn is faster than the third, all of which are faster than the main memory (Figure:7).

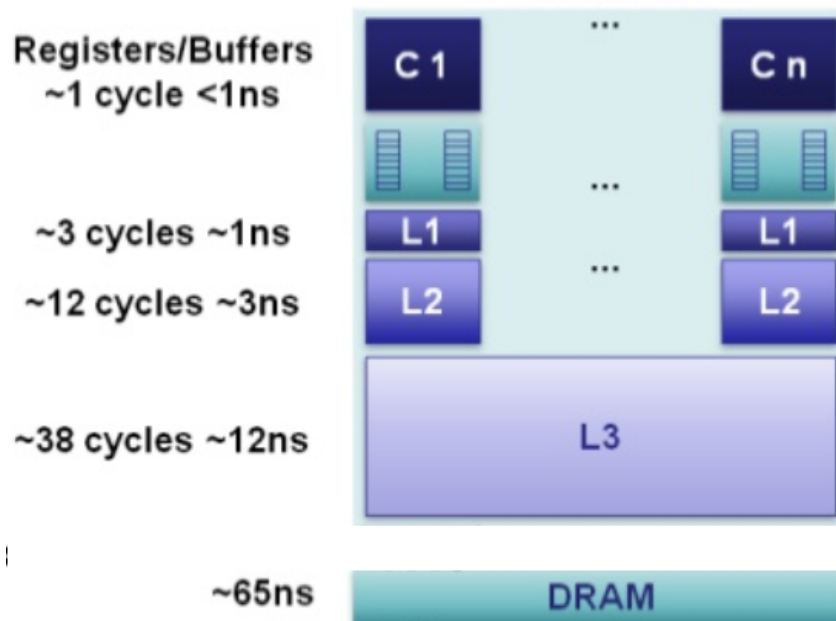


Figure 7: Speed of the different levels of cache and DRAM

When accessing some data from the memory, the L1 cache is first checked for and if it had been previously loaded there we can utilize the L1's fast access time to get the data, and we call that a **cache hit**, and when we don't find the wanted data in it we call it a **cache miss** instead.

If there was a cache miss on the L1 cache we try the L2, and if that misses we try the L3 and then we finally try the main memory (DRAM) if even that misses.

5.1.1 DRAM vs SRAM

The dynamic random-access memory (DRAM) and static random-access memory (SRAM) cells differ on an architectural level (Figure:8), where each cell of the SRAM is constructed of 6 cleverly placed transistors. Whereas the DRAM uses a single transistor and a capacitor, the capacitor in where the actual bit of data is stored which is the cause of slow access time since the capacitor needs to be completely depleted in order to read the bit value stored, and considering the size of the capacitors here the capacitor discharges automatically with time, which adds the requirement for a constant refresh of the cells to recharge the cells, this usually happens every 64ms and is also part of the reason the DRAM is slower. [4]

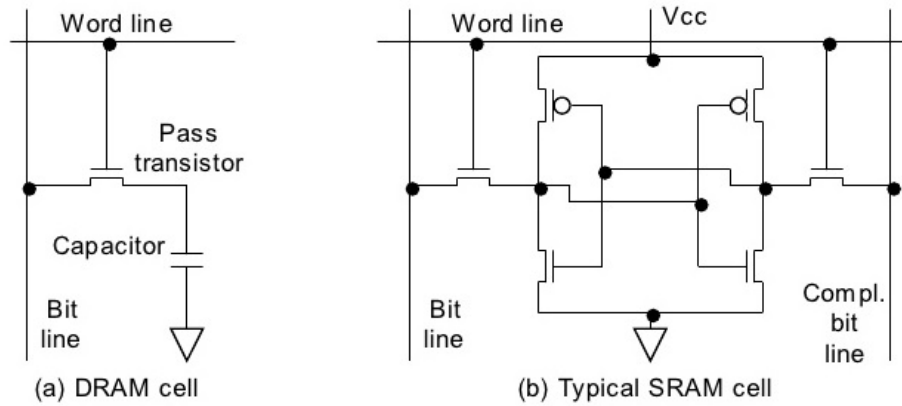


Figure 8: Cell architecture of DRAM and SRAM

5.1.2 Cache lines

Cache lines are the smallest block of memory that we can access *i.e.* load into the cache, they are usually 64 bytes or 32 bytes. The idea is that if we even load a single byte from the memory, the entire cache line that byte belongs to is loaded into the cache. Meaning if we try to access the byte right after it in memory (which is most probably the case) we can simply retrieve it from the cache instead.

The cache line is the same for all different levels of the cache, what changes however is of course how fast the memory level gets filled, the first level having the least amount of storage usually fills up first and so we have less L1 hits than L2 or L3 overall.

- X86 CPUs
- 64 bytes
- ARM CPUs
- 32 bytes

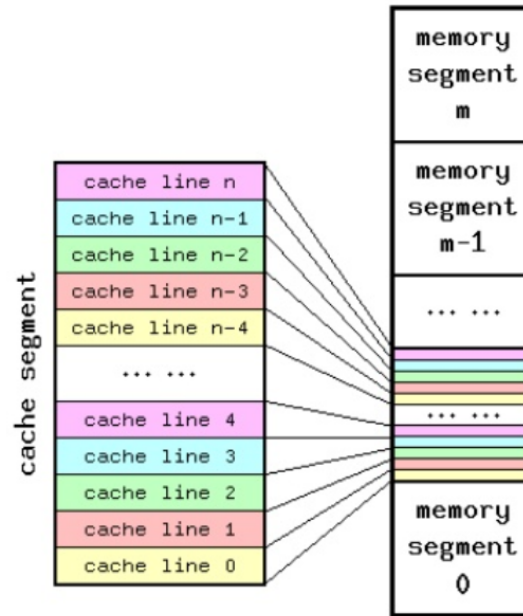


Figure 9: Memory cache lines

5.1.3 The cause

From we can already see that the use of an array of the (type.coeff) with a size of 56 bytes would fill almost the entirety of the cache line 10, meaning overall we'll have almost no cache hits, even if we access the coefficients array elements contigously (*ex*: coeff[0] then coeff[1] then coeff[2] and so on).

Hit rate $\sim 3/18 \sim 16.6\%$

Hit rate $\sim 7/8 \sim 87.5\%$

With an array of coeff_vals

With an array of floats (4 bytes)

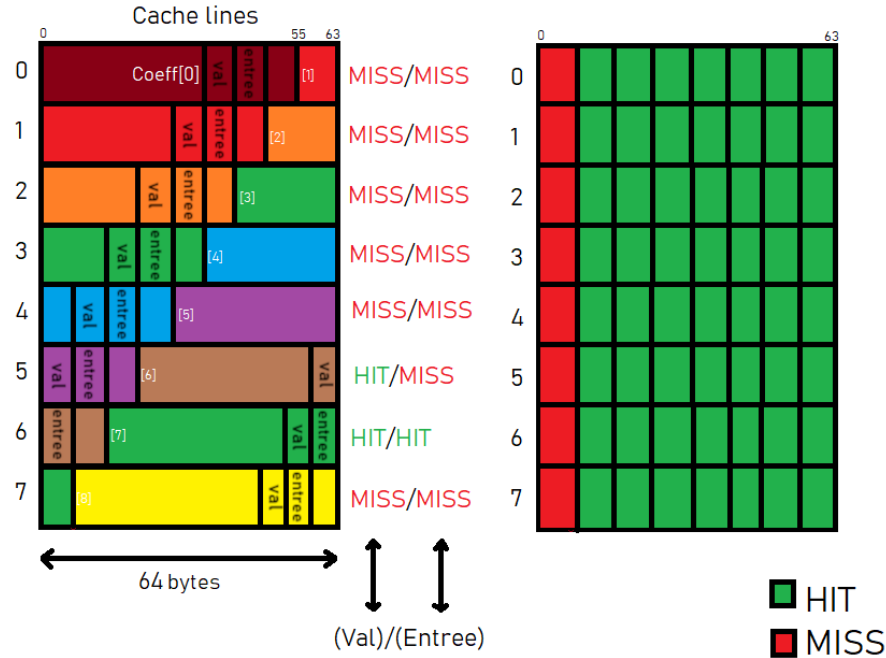


Figure 10: Cache hits on an array of type_coeff and another of floats

All of this would mean we would barely be using the cache (16.6% hit rate) and the computations will be dependent on the main memory's DRAM. which completely throttles the CPUs calculations since a since an access to the DRAM takes tens of clock cycles (Figure:8).

This here constitutes a bottleneck that has to be resolved before trying to parallelize the code, and it also explains the insignificant improvement from parallelizing the code previously where there was only a 20% performance increase on a 4 core machine, where it should have been 4 times faster.

6 Solution and results

Ideally the best solution would be to turn all the fields of the type_coeff structure into arrays, and also we would also need to remove the s field that references the next coefficient.

OLD STRUCT:

```
typedef struct type_coeff {
    float val;
    float proba;
    float Nbre_ES;
    float Nbre_S;
    float Nbre_E;
    int entree;
    int type;
    int evolution;
    float moy;
    float smoy;
    int gpe_liaison;
    struct type_coeff *s;
} type_coeff;
```

NEW STRUCT:

```
typedef struct type_coeff {
    float *val;
    float *proba;
    float *Nbre_ES;
    float *Nbre_S;
    float *Nbre_E;
    int *entree;
    int *type;
    int *evolution;
    float *moy;
    float *smoy;
    int *gpe_liaison;
} type_coeff;
```

And we would also need to change the type_neuron structure to have a single element of type_coeff (instead of a pointer):

OLD STRUCT:

```
typedef struct type_neuron {
    float seuil;
    int flag;
    float last_activation;
    float cste;
    int groupe;
    int nbre_voie;
    int nbre_coeff;
    int nbre_coeff_neuromod;
    char max;
    float posx;
    float posy;
    float posz;
    type_coeff *coeff; ---->
    type_coeff *nsor; ---->
} type_neuron;
```

NEW STRUCT:

```
typedef struct type_neuron {
    float seuil;
    int flag;
    float last_activation;
    float cste;
    int groupe;
    int nbre_voie;
    int nbre_coeff;
    int nbre_coeff_neuromod;
    char max;
    float posx;
    float posy;
    float posz;
    type_coeff coeff;
    type_coeff nsor;
} type_neuron;
```

This of course is a major architectural change and would require a lot of other changes and adjustments to the code base, especially the deletion of the s field from the type_coeff.

So instead, for now, in order to benchmark the impact of a similar change on the code, two float matrices for the 'val' and 'entree' (named respectively coeff_vals and coeff_entrees) fields of the coefficients were used instead, and were initialized and updated with the same values as the actual values from the type_coeff structure. The changes are limited to the hebbian neural network (prom.user/src/NN_Core/classique.rn/trad_neuron.c) and the kohonen self-organizing

map (same directory in kohonen.c). and the results seem to reflect that the changes solved the bottleneck and the code can run in parallel (Figure:11) (Figure:12).

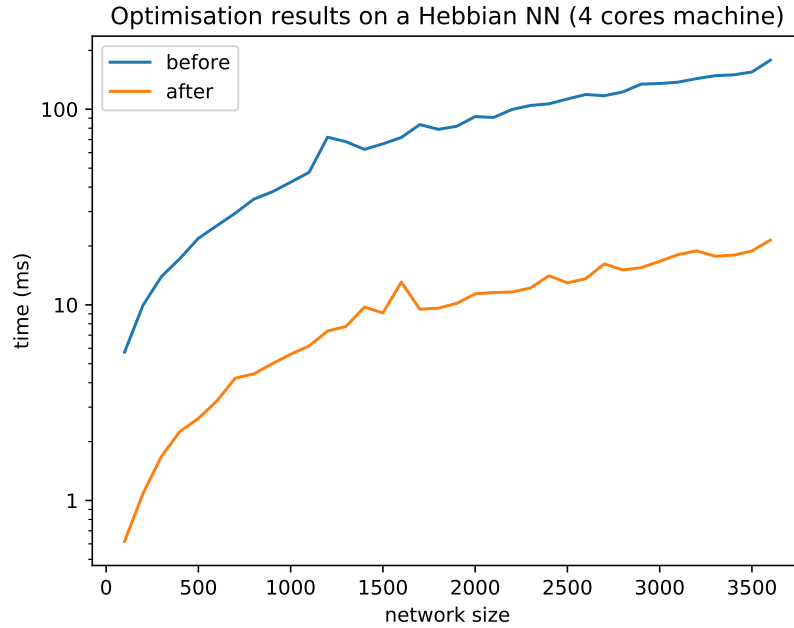


Figure 11: Results of the optimisation on hebbian NN

6.1 tests

To verify that the program continues to function correctly some tests were created that compare results (coefficient values) of the original version of the program (without the optimisation) with the new version.

In order to do this the two versions of the program were used to compile a test .script file with their cc_letto (used to generate a .res file used when executing the promethe simulator) that prompts for a seed for the random number generator, these programs when launched (using promethe) generate a .SAVE file that contains a text version of the coefficients and their values. The idea of the test is that given the same random seed, these two .SAVE files should be the same for the two versions.

For now these tests still need some more work done on them since they are not working. There is still a lot of work to be done here in order to utilize the new optimisation and there's the question of how much more *should* be done,

since these are major changes to a very large code base, and would therefore require a lot of time to implement.

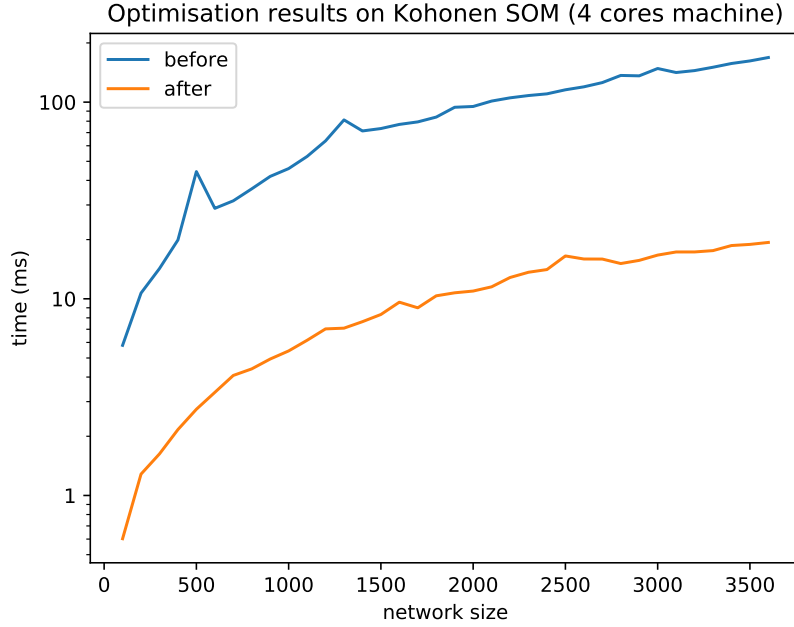


Figure 12: Results of the optimisation on kohonen SOM

7 Conclusion

It seems that in order to utilize the cache's fast access time, and to avoid memory access bottlenecks. The data needs to be divided into contiguous arrays of small chunks (few bytes, float for example) of the same field (for example 'val') this makes sure that cache lines will contain more of these small chunks as opposed to having an array of a big structure (tens of bytes) which would cause the cache lines to only one or two copies of the structure, plummeting the hit rate on the cache.

This is a result of the fact that we only use a few fields of the structure at any given time (we only use the 'val' and 'entree' fields when updating hebbian NN neurons), so most of the cache line is not used and wasted. However when there are multiple arrays of each field (floats) or small collection of fields, we would have more elements of the same type in our cache lines and therefore more hits of accelerated memory access of the CPU's cache.

References

- [1] M. E. Wolf and M. Lam, “A data locality optimizing algorithm,” *Computer Systems Laboratory Stanford University, CA 94305*, 1991.
- [2] M. Wolfe, “More iteration space tiling,” *Oregon Graduate Center*, 1989.
- [3] J. Xue, *Loop Tiling for Parallelism*. Kluwer Academic Publishers, 2000.
- [4] U. Drepper, “What every programmer should know about memory,” *Red Hat, Inc.*, 2007.